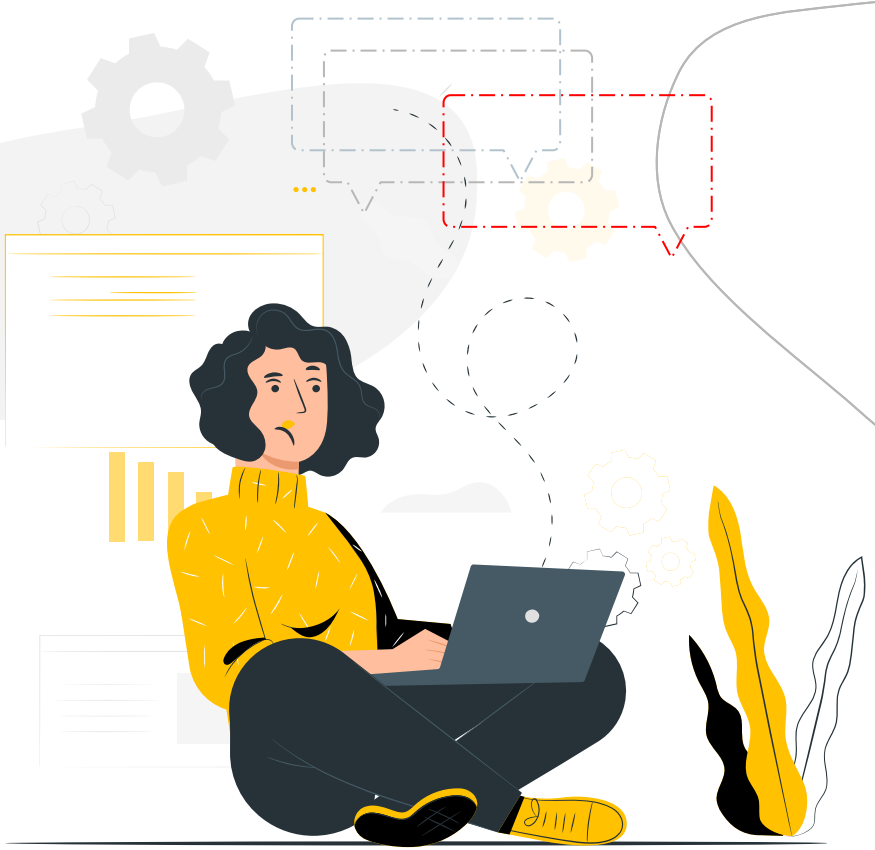

Discussion Detox

Multilingual Machine Learning algorithms to identify toxic comments on the internet

Author: Drenizë Rama

TABLE OF CONTENTS

**“INTERNET RULE #1:
Never read the comments.”**

— WIRED

02 | the data

Example comment #1:

"What a motherfucking piece of crap those fuckheads for blocking us!"

Example comment #2:

"Hey, faggot.
You fucking retard. You better quit undoing my vandalism, bitchboy."

Example comment #3:

"but ew
He was a fag which is against nature and is the most disgusting thing. Youre not a woman are you? Sexism is wrong. Being wrong is for women."

An **online newspaper** or a **social media web host** wants to keep the discussions under each article clean and respectful

However, going through every comment manually is tiresome and very expensive

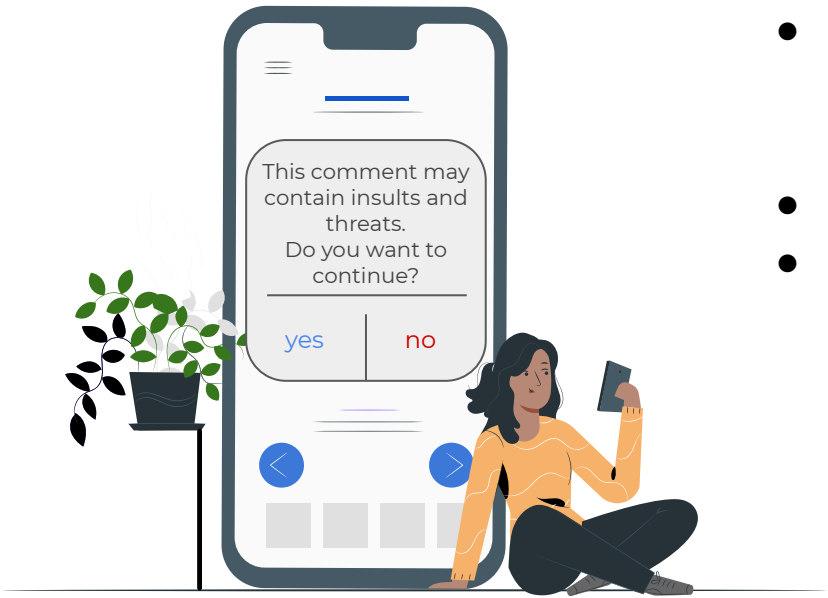

goal

- build a **natural language algorithm** that classifies social media comments into toxic and non-toxic categories
- at a **low cost**
- across **different languages**

02 | the data

The overall amount of toxic comments is 11.4%

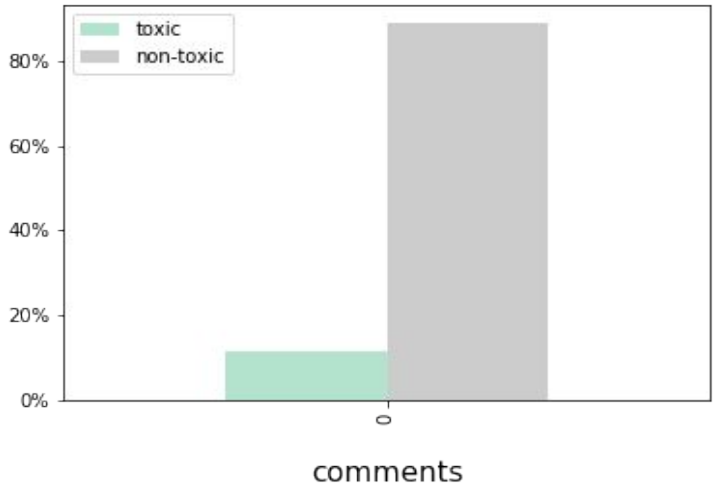

- Data was **provided by Google** and Jigsaw
- **Publication** dates of the comments range from **2015 to 2017**
- **223,549** comments in train set

02 | the data

The more hate - the shorter the comments

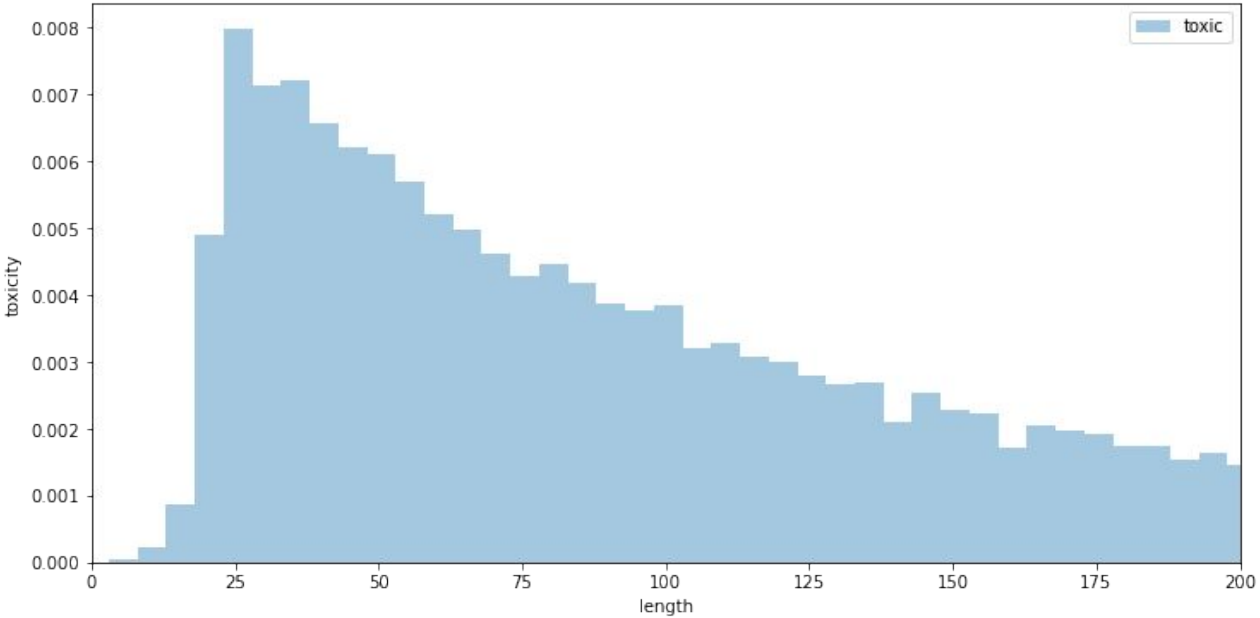

frequently used words

In toxic comments

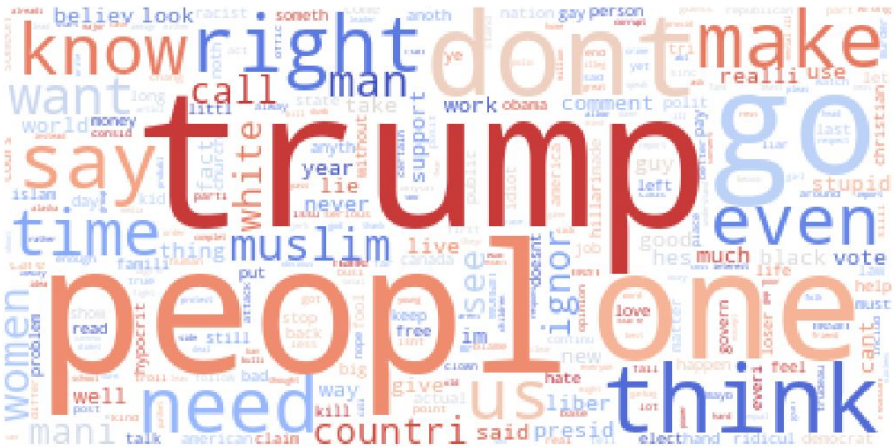

In comments containing identity attack

03 | methods

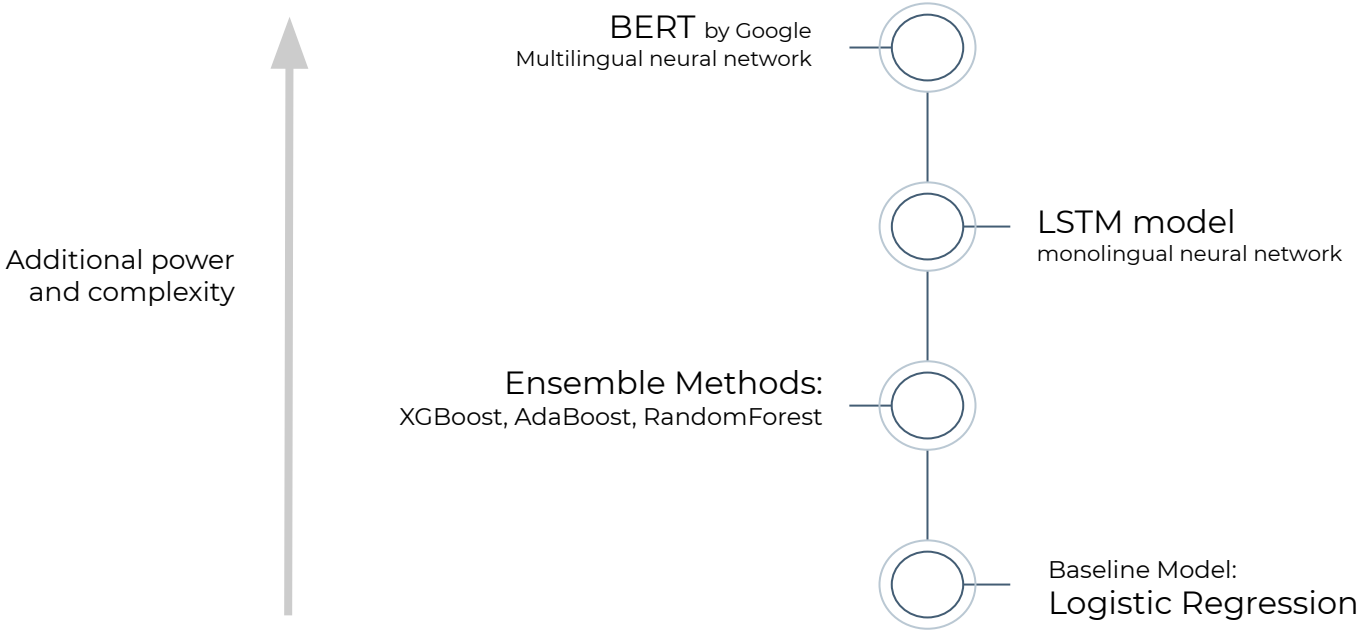
Additional power
and complexity
BERT by Google
Multilingual neural network
LSTM model
monolingual neural network
Ensemble Methods:
XGBoost, AdaBoost, RandomForest
Baseline Model:
Logistic Regression

How Natural Language Processing works

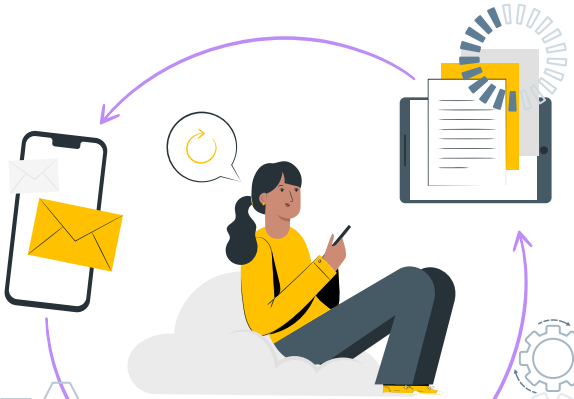

Text Cleaning
Stemming, remove symbols, emojis, punctuation, stopwords

Train Model
Apply the model in data

Build Model
Adjust type, layers, epochs etc.

Encode
Create vectors for word representation

Tokenize
Divide sentences into sequences of words or characters

At what rate do the models predict toxic comments correctly? (recall)

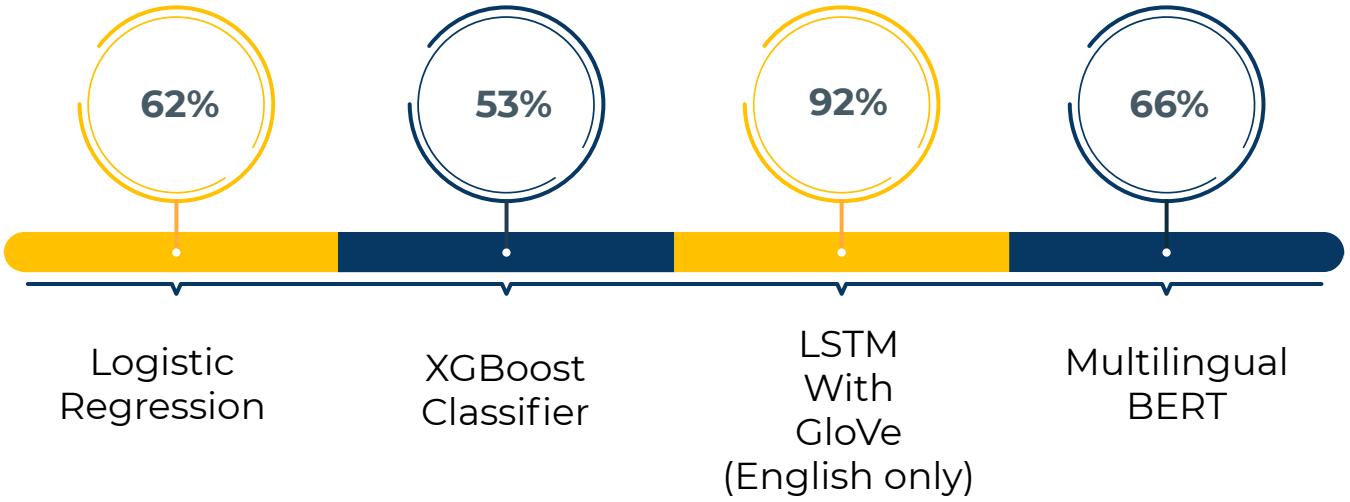

Multilingual Toxic Comments Classifier on Google Cloud Platform

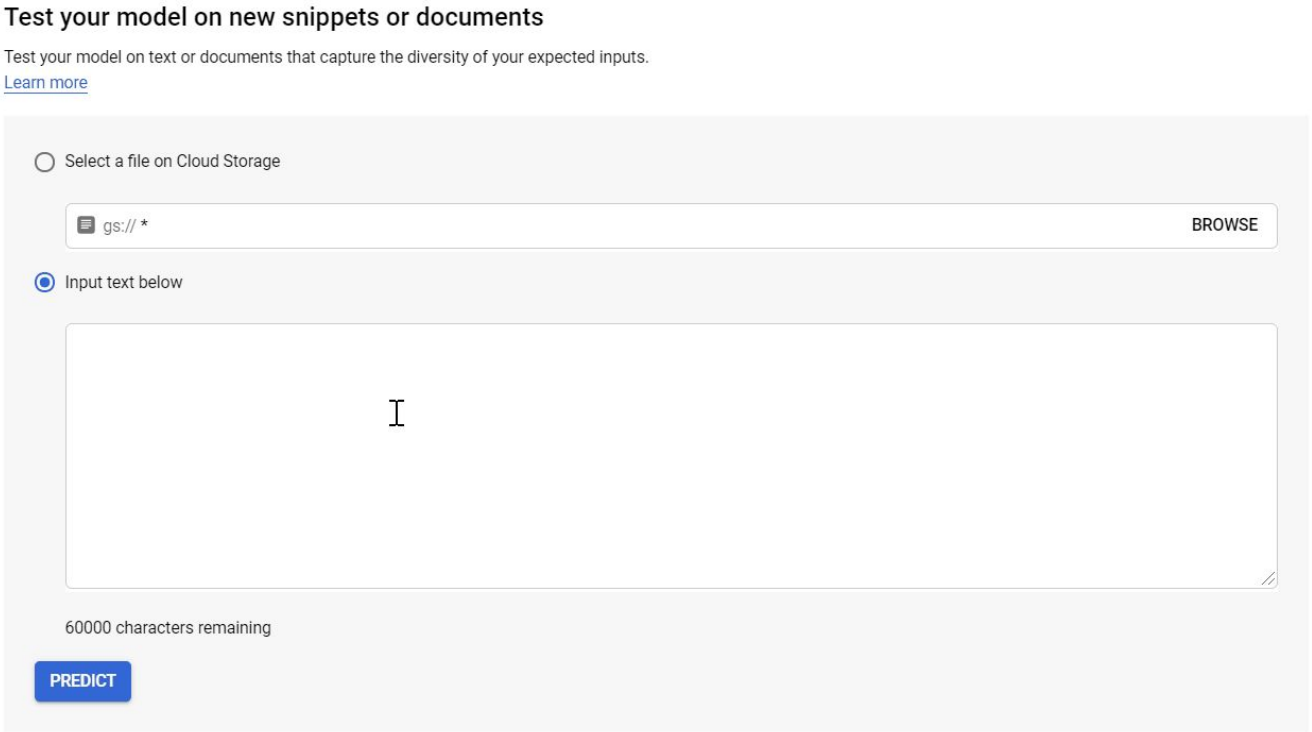

Appendix

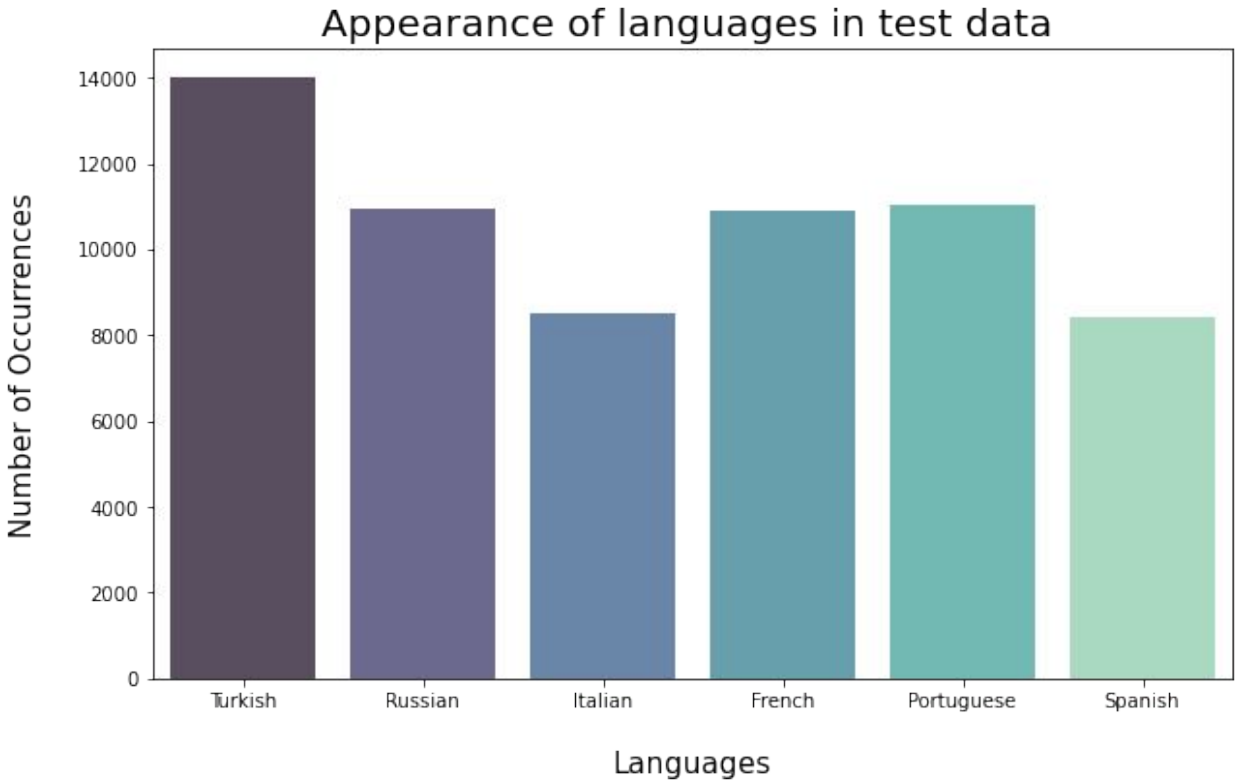
Appearance of languages in test data
Number of Occurrences
14000
12000
10000
8000
6000
4000
2000
0
Turkish
Russian
Italian
French
Portuguese
Spanish
Languages

05 | recommendations

- reduce costs of identifying toxic comments with simple models
- Use pre-trained word embeddings to improve model performance
- Invest in training multilingual models to secure competitiveness in the future

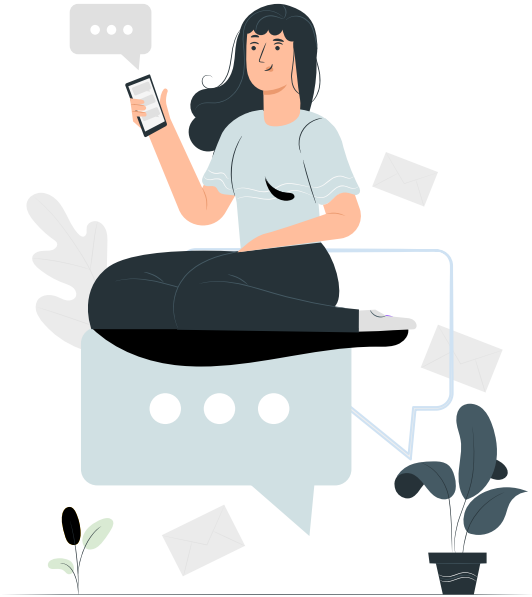

06 | future work

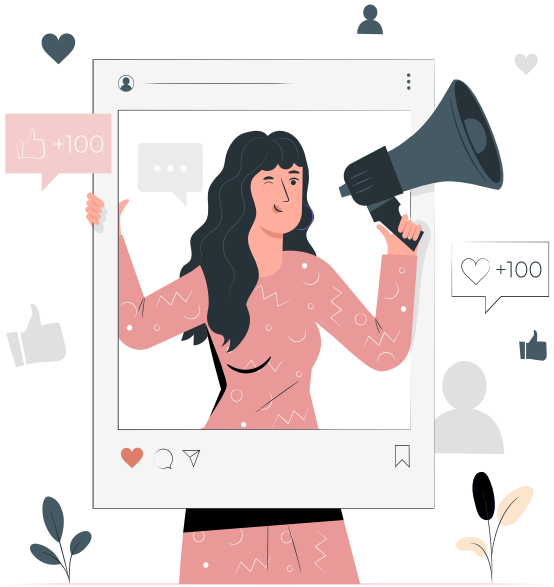

- Optimize model to reduce bias and recognize irony and implicit aggression
- Create a web tool that recommends users to adjust their language before posting a comment
- time-series analysis of toxicity online
- Work with new tools like LASER by Facebook

NLP throughout different languages

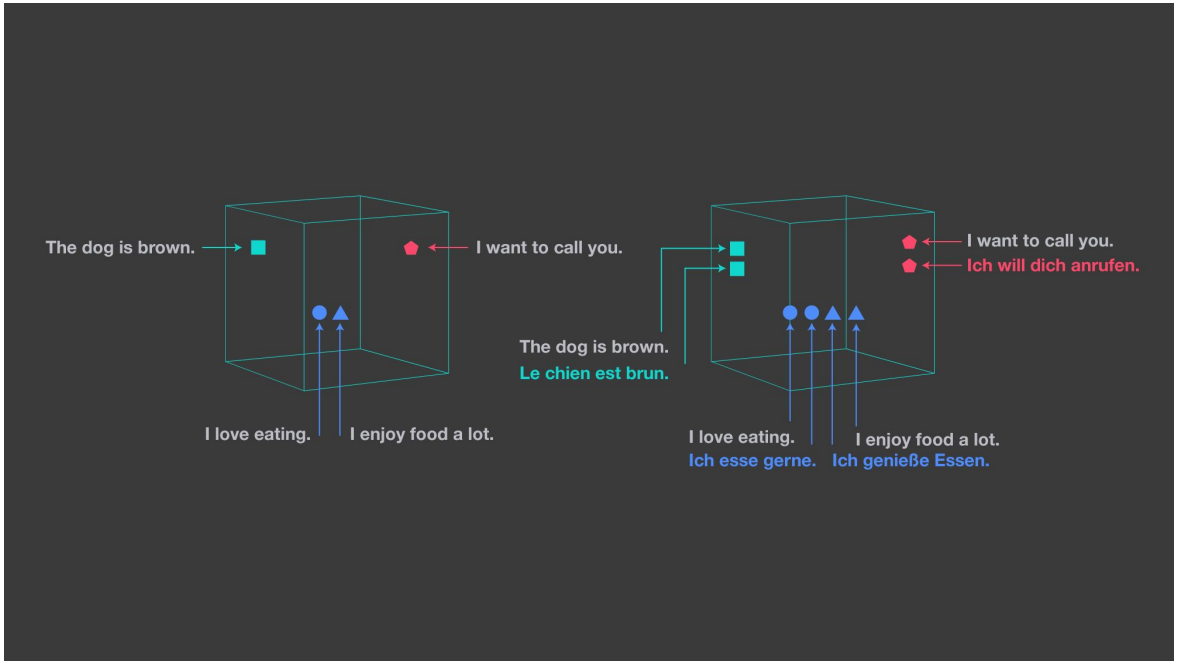

06 | future work

protect users from toxic content

Approach 1

Approach 2

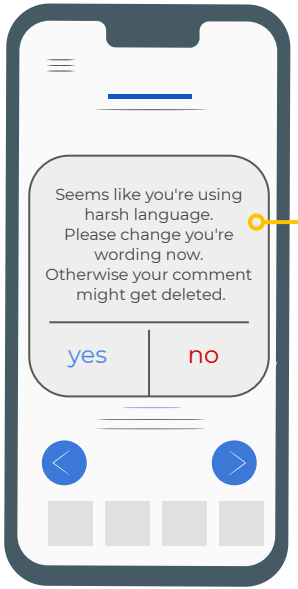

mirror toxic behavior in advance

Ja, und?

VICE Artikel Videos Politik Tech Menschen Drogen Popkultur

Tech

Stiftung Warentest veröffentlicht Geldstrafen-Katalog für Hasskommentare

„Merkel muss öffentlich gesteinigt werden" kostet 2000 Euro.

Von Johannes Hausen

19 Mai 2016, 12:03pm Teilen Twittern Snap

Thank you

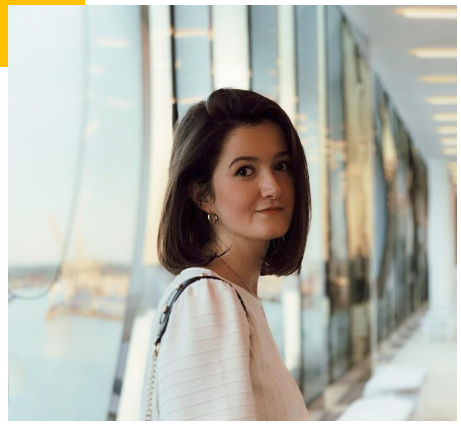

Drenizë Rama

Data Scientist

https://drenize.github.io/

https://www.linkedin.com/in/dreniz%C3%AB-rama-6121a4157/

Appendix

LSTM architecture

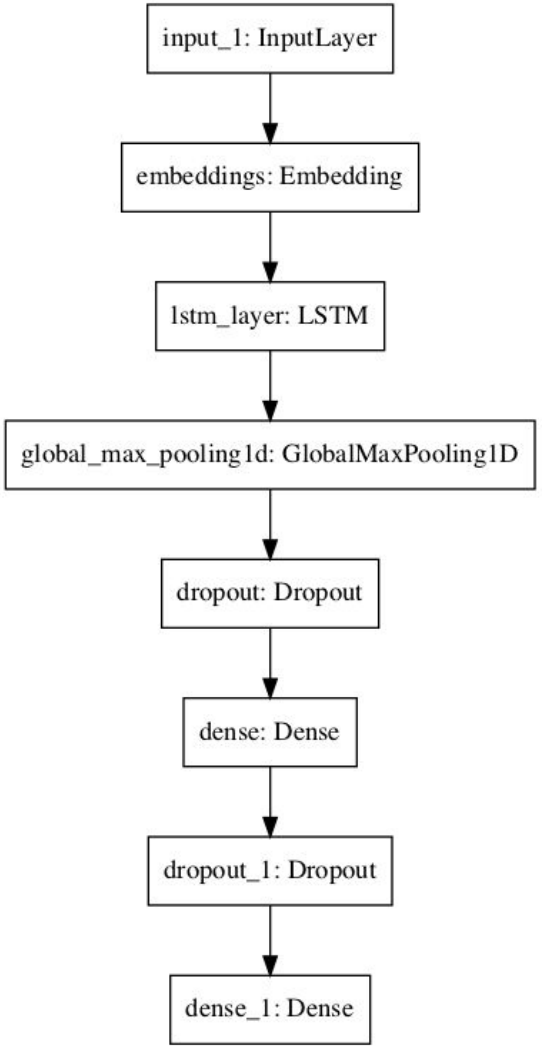

Appendix

BERT architecture

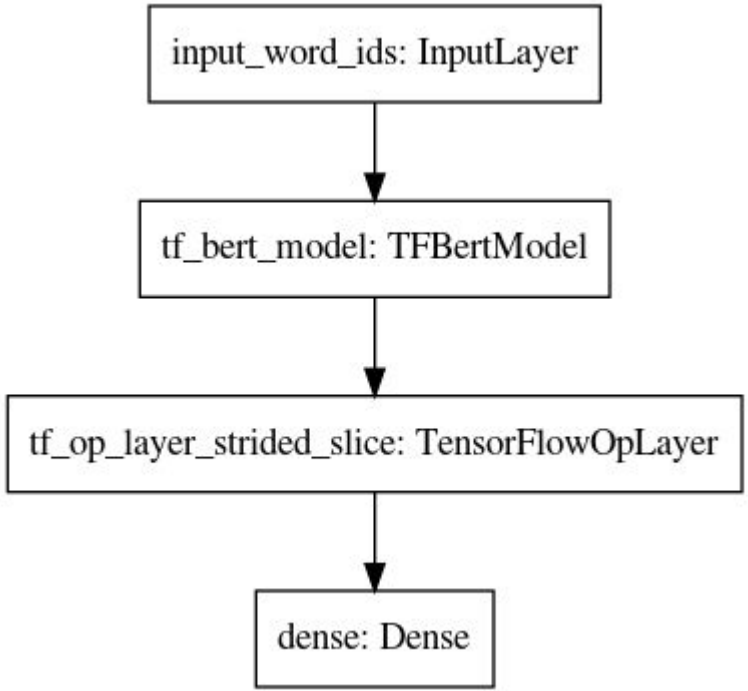

Appendix

Baseline Model: Logistic Regression

```
[[49808   853]
 [ 1991  3236]]
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	50661
1	0.79	0.62	0.69	5227
accuracy			0.95	55888
macro avg	0.88	0.80	0.83	55888
weighted avg	0.95	0.95	0.95	55888

Appendix

XGBoost Classifier

```
Confusion Matrix :
[[149624   2007]
 [  7607   8423]]
Accuracy Score : 0.9426581017648707
Report :
              precision    recall  f1-score   support

           0       0.95      0.99      0.97    151631
           1       0.81      0.53      0.64     16030

    accuracy                           0.94    167661
   macro avg       0.88      0.76      0.80    167661
weighted avg       0.94      0.94      0.94    167661
```

Appendix

LSTM model with GloVe

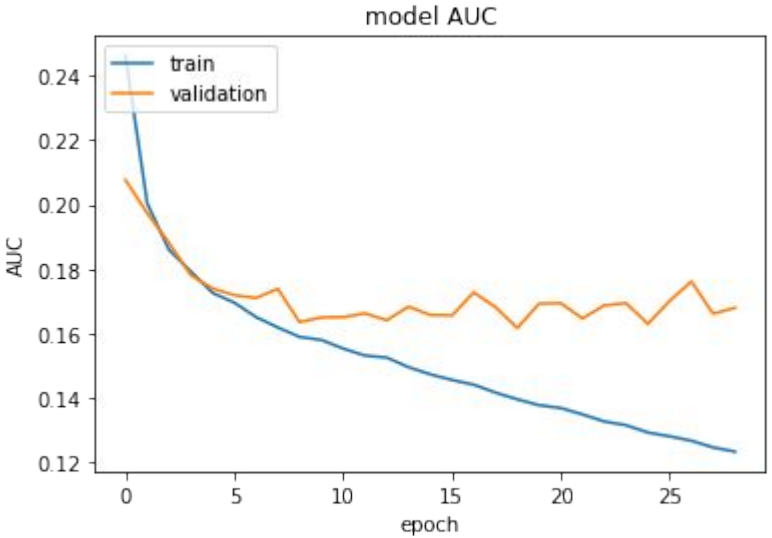

Appendix

BERT multilingual model

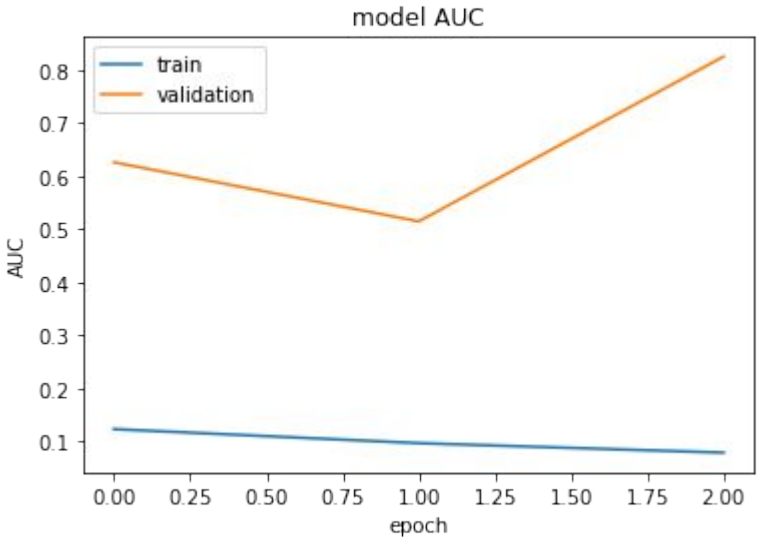

Appendix

How does a Neural Network work?

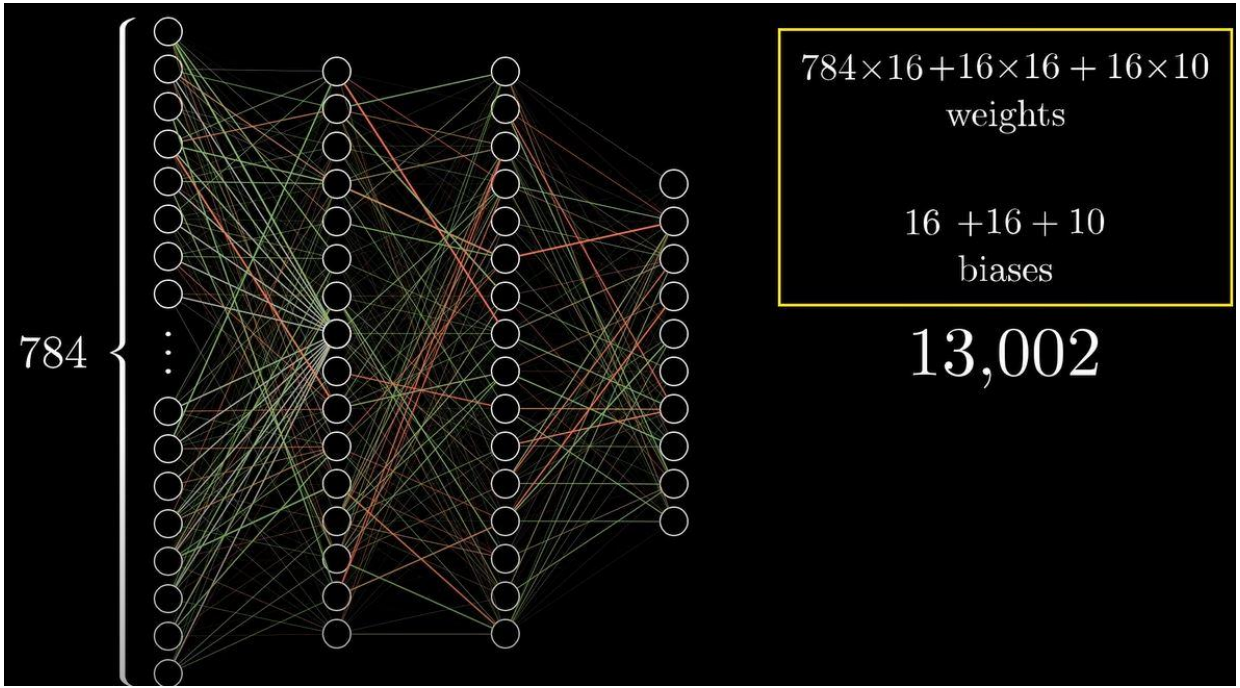

Appendix

GloVe: Global Verctors for Word Representation

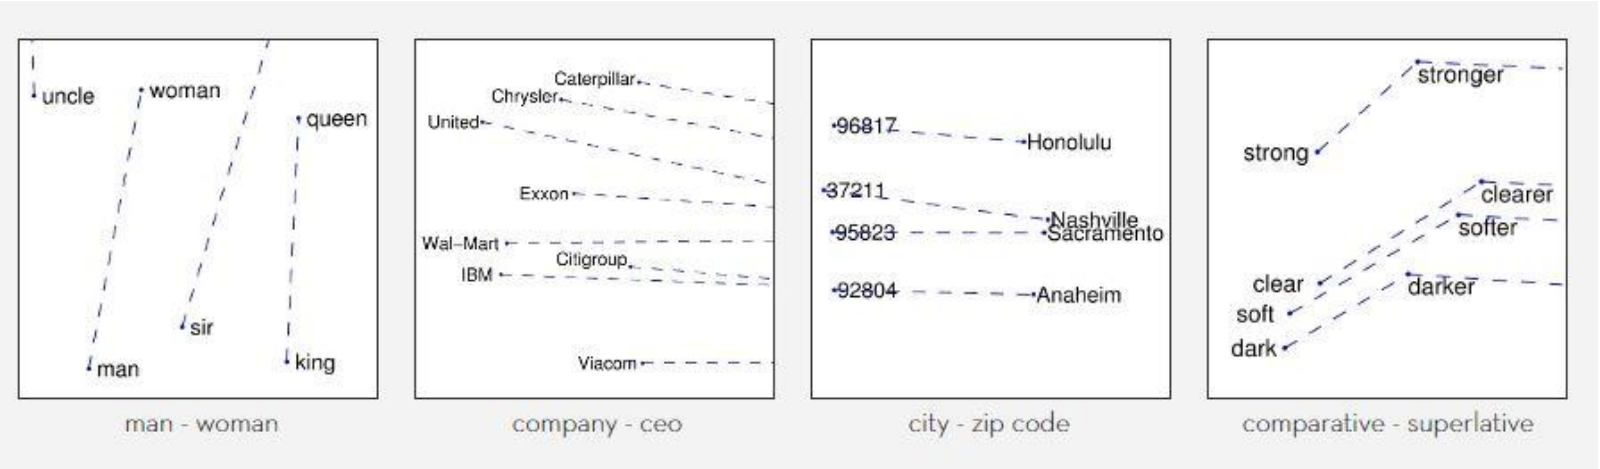

Hasskommentare
im Netz
identifizieren

Discussion Detox

App

by Drenizë Rama